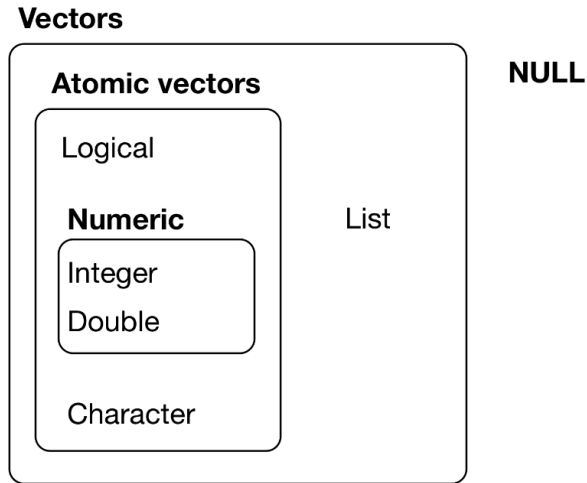


Data Manipulation with dplyr

36-600

Lists



- As can be seen above, a list is *not* an atomic entity; it may be thought of as a (perhaps heterogeneous) collection of atomic vectors:

```
(x <- list(1:5,c("a","b")))
```

```
## [[1]]  
## [1] 1 2 3 4 5  
##  
## [[2]]  
## [1] "a" "b"
```

Naming List "Columns"

- When you initialize a list, you can name each of the vectors and then use the names to access list elements

```
x <- list(u=1:5,v=c("a","b"))  
x
```

```
## $u  
## [1] 1 2 3 4 5  
##  
## $v  
## [1] "a" "b"
```

```
x$v
```

```
## [1] "a" "b"
```

```
x$u[3]
```

```
## [1] 3
```

- Note the use of the dollar sign...the name of the list goes to the left of the dollar sign, while the name of the list element goes to the right

Data Frames

- A data frame is simply a list whose entries all have the same number of elements!

```
(x <- data.frame(u=1:2,v=c("a","b"),w=c(TRUE,FALSE)))
```

```
##   u v      w  
## 1 1 a   TRUE  
## 2 2 b  FALSE
```

- Note that when we output a data frame, row numbers are added
- Are we sure this is a list?...we can check the variable type:

```
typeof(x)
```

```
## [1] "list"
```

```
class(x)
```

```
## [1] "data.frame"
```

Dataset: Galaxy Properties

- Let's read in data that are stored on disk in an ASCII (i.e., human-readable) file in csv format:

```
df <- read.csv("http://www.stat.cmu.edu/~pfreeman/GalaxyMass.csv", stringsAsFactors=TRUE)
```

- df stands for "data frame," but you can use whatever variable name you wish
- The data contain 3456 galaxies (in rows), with 10 measurements for each (in columns):

```
dim(df)
```

```
## [1] 3456 10
```

- We can use the head() function to display the first (in this case, 2) rows of the data frame:

```
head(df, 2)
```

```
##      field      Gini      M20      C      A      size      n      q z.mode
## 1 COSMOS 0.4132853 -1.132330 2.257530 0.1211518 0.7554244 0.3398 0.3579 2.04
## 2 COSMOS 0.4177935 -1.603147 2.970555 0.1002479 0.6537786 0.6722 0.8124 1.79
##      mass
## 1 10.55120
## 2 9.87608
```

Dataset: Galaxy Properties

- To see the names associated with each variable/column, use the `names()` function:

```
names(df)
```

```
## [1] "field" "Gini" "M20" "C" "A" "size" "n" "q"  
## [9] "z.mode" "mass"
```

- The first column represents a sector of the sky in which the galaxy lies
 - this is a *factor* variable, because it takes on a few discrete values that don't have a specific ordering

```
unique(df$field)
```

```
## [1] COSMOS EGS GOODSN GOODSS UDS  
## Levels: COSMOS EGS GOODSN GOODSS UDS
```

dplyr

- dplyr is a package within the larger tidyverse that one uses to manipulate data frames

```
suppressMessages(library(tidyverse)) # suppressMessages() makes for cleaner output
```

- It may be helpful to think of data frames as nouns and dplyr functions as verbs, actions that you apply to the data frames
- The following are the most basic dplyr verbs:
 - slice(): choose particular rows based on integer indexing
 - filter(): choose particular rows based on logical criteria
 - group_by(): split the data frame into groups of rows based on factor variable values
 - select(): choose particular columns
 - arrange(): order rows by value of a column
 - mutate(): create new columns
- **NOTE:** calling dplyr verbs always outputs a new data frame...it *does not alter* the existing data frame

The slice() Function

- Use the `slice()` function when you want to retain certain rows:

```
df %>% slice(.,c(7,14)) # output rows 7 and 14 of the data frame
```

```
##      field      Gini      M20      C      A      size      n      q z.mode
## 1 COSMOS 0.3319734 -1.29408 2.975167 0.09361421 0.8088991 0.7442 0.2343 1.96
## 2 COSMOS 0.4839426 -1.60984 3.005040 0.17674453 0.4727087 1.3563 0.4967 1.69
##      mass
## 1 10.6824
## 2 10.1826
```

- The `%>%` is a *pipe*
 - a pipe takes the output of one R command and uses it as input to a following command
 - here, we pipe the output of `df` (which is simply the whole data frame) to the `slice()` function
 - the period in the `slice()` function call shows where the data frame is supposed to go among the functional arguments
- Slicing can be done "negatively":

```
df %>% slice(.,-c(1:2,19:23)) %>% nrow(.) # how many rows are left after removing rows 1-2 and 19-23?
```

```
## [1] 3449
```


The filter() Function

- Use the filter() function when you want to choose rows based on logical conditions:

```
df %>% filter(.,field=="COSMOS") %>% head(.,2)
```

```
##      field      Gini      M20      C      A      size      n      q z.mode
## 1 COSMOS 0.4132853 -1.132330 2.257530 0.1211518 0.7554244 0.3398 0.3579 2.04
## 2 COSMOS 0.4177935 -1.603147 2.970555 0.1002479 0.6537786 0.6722 0.8124 1.79
##      mass
## 1 10.55120
## 2  9.87608
```

```
df %>% filter(.,mass>10) %>% nrow(.)
```

```
## [1] 1862
```

```
df %>% filter(.,(field=="GOODSS" & mass<10)) %>% nrow(.)
```

```
## [1] 276
```

- Note the use of & in the third example: this combines conditions via the logical and (where | would be the analogous symbol for or)

The group_by() Function

- Use the group_by() function to split a data frame into groups of rows based on the values of a factor variable
 - note that group_by() in and of itself is only useful when its output is piped its to another function, e.g., summarize():

```
df %>% group_by(.,field) %>% summarize(.,Mean=mean(mass),Number=n())
```

```
## # A tibble: 5 × 3
##   field    Mean Number
##   <fct>  <dbl>   <int>
## 1 COSMOS  10.2     905
## 2 EGS     10.0     750
## 3 GOODSN  10.1     464
## 4 GOODSS  10.1     588
## 5 UDS     10.2     749
```

- Mean and Number are column names in the *output* data frame
 - what is shown are the average value of the galaxy masses in each field, along with the number of galaxies in each field
- Note that the output of summarize() is a "tibble", which is a tidyverse-specific alternative to a base-R data frame

The select() Function

- Use the select() function when you want to choose certain columns:

```
df %>% select(.,Gini,q,z.mode) %>% slice(.,c(1:4)) # here we choose columns *and* rows
```

```
##           Gini          q z.mode
## 1 0.4132853 0.3579    2.04
## 2 0.4177935 0.8124    1.79
## 3 0.4212831 0.3585    2.02
## 4 0.4725081 0.4927    1.94
```

- select() can also be used negatively:

```
df %>% select(.,-Gini,-q,-z.mode) %>% ncol(.) # how many variables/columns are left?
```

```
## [1] 7
```

The arrange() Function

- Use the arrange() function to order rows by values of a column:

```
df %>% arrange(.,desc(mass)) %>% select(.,C,A,z.mode) %>% head(.,2)
```

```
##           C           A z.mode
## 1 3.594274 0.1711388    1.68
## 2 3.655988 0.2116704    1.78
```

- If one uses arrange(.,mass) then the masses will be ordered in *ascending* order, by default

The mutate() Function

- Use the mutate() function when you want to create one or several columns:

```
df %>% mutate(.,mass.linear=10^mass) %>% select(.,mass,mass.linear) %>% arrange(.,mass) %>% head(.,3)
```

```
##      mass mass.linear  
## 1 8.285    192752491  
## 2 8.540    346736850  
## 3 8.550    354813389
```

Saving the New Data Frame

- As stated above, the original data frame is not altered in piping operations
- To save the result of a series of piping operations, use `<-` or `->`:

```
df.new <- df %>% filter(.,field=="EGS") %>% select(.,z.mode,mass)
nrow(df.new)
```

```
## [1] 750
```

```
df %>% filter(.,field=="COSMOS") %>% select(.,Gini,M20,C,A) -> df.new
nrow(df.new)
```

```
## [1] 905
```