

Pure Prediction: KNN and SVM

36-600

KNN: K Nearest Neighbors

- In words: KNN examines the k data points closest to a location \mathbf{x} and uses just those data to generate predictions
- KNN straddles the boundary between fully parameterized models like linear regression and fully data-driven models like random forest: a KNN model is data-driven, but one *can* actually write down a compact parametric form for it *a priori*
 - for regression:

$$\hat{Y}|\mathbf{x} = \frac{1}{k} \sum_{i=1}^k Y_i$$

- for classification:

$$P[Y = j|\mathbf{x}] = \frac{1}{k} \sum_{i=1}^k \mathbb{I}(Y_i = j)$$

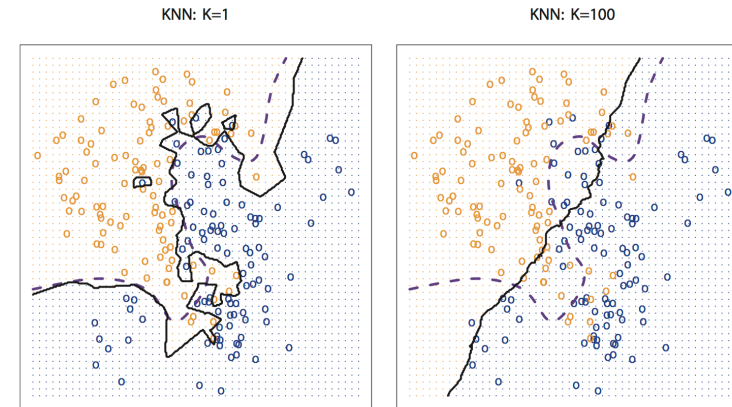
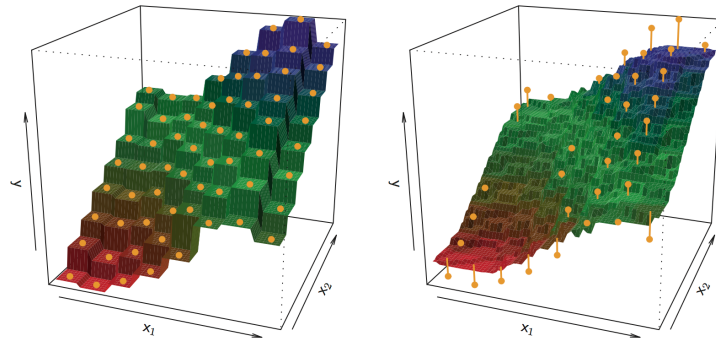
- the summations are over the points $\mathbf{x}_1, \dots, \mathbf{x}_k$ that are "closest" to \mathbf{x} (usually in a Euclidean sense)
- $\mathbb{I}(\cdot)$ is the indicator function: it returns 0 if the argument is false, and 1 otherwise

Finding the Optimal Number of Neighbors

- For KNN, the number of neighbors k is a *tuning parameter*
 - this means that in theory, in addition to splitting data into training and test datasets, *the training data themselves have to be split into a smaller training set and a validation set...*for each value of k , you train on the smaller training set, and compute MSE or MCR using the validation set
- Once the optimum value of k is determined, you re-run the model using the *full* (unsplit) training set, and assess the model using the test data
- The FNN package in R does cross-validation on the training set automatically, under the hood...we do not need to code cross-validation ourselves

Finding the Optimal Number of Neighbors

- Note that if...
 - k is too small, the resulting model is *too flexible*, which means it has low bias (it is right on average) but high variance (the predictions are more uncertain)
 - k is too large, the resulting model is *not flexible enough*, resulting in high bias (wrong on average) and low variance (nearly the same predictions, every time)



(Figures 3.16 and 2.16, *ISLR 1e*)

KNN in Context

- *As a general rule, parametric methods [like linear regression] will tend to outperform non-parametric approaches [like KNN] when there is a small number of observations per predictor.*
 - this is the *curse of dimensionality*: for data-driven models, the amount of data you need to get similar model performance goes up exponentially with p , the number of predictor variables
 - KNN might not be a good model to learn when the number of predictor variables is large
- *Even in problems in which the dimension is small, we might prefer linear regression to KNN from an interpretability standpoint. If the test MSE of KNN is only slightly lower than that of linear regression, we might be willing to forego a little bit of prediction accuracy for the sake of a simple model...*
 - KNN is not the best model to learn if inference is the goal of an analysis

KNN: Three Critical Points to Remember

- To determine which neighbors are the nearest neighbors, pairwise (Euclidean) distances are computed...so we need to scale (or standardize) the individual predictor variables
- Since it utilizes distances between data, *KNN only works with quantitative predictor variables*
- Don't blindly compute a pairwise distance matrix...for instance, if $n = 100,000$, then your pairwise distance matrix will have 10^{10} elements, each of which uses 8 bytes in memory...resulting in a memory usage of 80 GB!
 - subsample your data, limiting n to be $\lesssim 15,000$ -20,000
 - use a variant of KNN that works with sparse matrices (matrices that can be compressed since most values are zero)
 - make use of a "kd tree" to more effectively (but only approximately) identify nearest neighbors
 - find a computer system with 32 GB or 64 GB (or more) of memory
 - the FNN package in R has an option to search for neighbors via the use of a kd tree; apply this option if you have more than, e.g., 10,000 rows in your data frame

KNN: Regression Example

- We input and split a 10,000-galaxy dataset that we will use to predict distances of galaxies from the Earth
 - note that we scale the predictor variables (but not the response), using the `scale()` function

```
library(FNN)

k.max <- 20
mse.k <- rep(NA,k.max)
for ( kk in 1:k.max ) {
  knn.out <- knn.reg(train=df.train[,1:6],y=df.train[,7],k=kk,algorithm="brute")
  mse.k[kk] <- mean((knn.out$pred-df.train[,7])^2)
}
k.min <- which.min(mse.k)
cat("The optimal number of nearest neighbors is ",k.min,"\n")
```


The optimal number of nearest neighbors is 11

- If the optimal number of nearest neighbors is k_{\max} , then your value of k_{\max} is too small...increase it and start over!

KNN: Regression Example

```
suppressMessages(library(tidyverse))

ggplot(data=data.frame("k"=1:k.max,"mse"=mse.k),
       mapping=aes(x=k,y=mse)) +
  geom_point() + geom_line() +
  xlab("Number of Nearest Neighbors k") +
  ylab("Validation MSE") +
  geom_vline(xintercept=k.min,color="red")
```




KNN: Regression Example

- Given `k.min`, we generate test-set predictions and generate the test-set MSE

```
knn.out <- knn.reg(train=df.train[,1:6],  
                  test=df.test[,1:6],  
                  y=df.train[,7],  
                  k=k.min,algorithm="brute")  
(knn.mse <- mean((knn.out$pred-df.test[,7])^2))
```

```
## [1] 0.07574588
```

```
ggplot(data=data.frame("x"=df.test[,7],  
                      "y"=knn.out$pred),  
       mapping=aes(x=x,y=y)) +  
  geom_point(size=0.1,color="saddlebrown") +  
  xlim(0,2) + ylim(0,2) +  
  xlab("Observed Test-Set Response") +  
  ylab("Predicted Test-Set Response") +  
  geom_abline(intercept=0,slope=1,color="red")
```



KNN: Prediction in Classification Models


- Note that for classification:
 - instead of using `knn.reg()`, you would use `knn.cv()` to determine the optimum value of k ; and
 - you would use `knn()` to generate predictions on the test set
- Note the `cl` argument for both: this is where `resp.train` goes (since `cl` means "class")...change `y` to `cl`
- `knn()` will output class predictions (based on majority vote!), unless you add the argument `prob=TRUE` to the call
 - if you want to compute, e.g., a ROC curve for a KNN classifier, use the (pseudo-)code below to extract the Class 1 probabilities

```
knn.out      <- knn(...,prob=TRUE)
knn.prob     <- attributes(knn.out)$prob
w            <- which(knn.out=="<class 0>") # insert name of Class 0 here
knn.prob[w]  <- 1 - knn.prob[w]    # knn.prob is now the Class 1 probability!
```

SVM: Support-Vector Machine

- A *support vector machine* is an enigmatically named machine learning algorithm most often used for classification
- In theory (but not in practice...because of algorithmic cleverness), SVM transforms predictor data into a *higher-dimensional space* and in that space constructs a linear boundary that optimally separates instances of two classes
 - SVM is not designed to tackle analyses in which the response has more than two classes!
- Like a KNN model, an SVM model is a purely predictive model, useless for inference...
- ...and like a KNN model, an SVM model utilizes distances between data, so scaling the predictor data is required...
- ...and like a KNN model, an SVM model works only with quantitative predictor variables
- The *ISLR* textbook gives a somewhat mathematical description of SVM, and you should read the SVM chapter (chapter 9 of the first edition) if you are interested in such details


But First: What *is* a Support Vector, Anyway?



(Figure 9.3, *ISLR 1e*)

- This figure shows an example of a *maximum margin classifier* that determines the linear boundary in the native space of the predictor data that has the largest *minimum* distance to a training datum
 - (MMC requires a complete separation between instances of the two classes, so it is really never used in practice)
 - however: look at the three short line segments that are perpendicular to the boundary...these are *support vectors*: they "hold up," or support, the linear boundary

Support Vector Machine



(Figure 9.9, *ISLR* 1e)

- SVM, as stated above, defines a linear boundary in a higher-dimensional space than the native space of the predictor variables
 - SVM with a polynomial kernel of degree 2 transforms a space where $p = 2$ to one where $p = 5$:
$$\{X_1, X_2\} \rightarrow \{X_1, X_2, X_1^2, X_2^2, X_1X_2\}$$
 - but it doesn't, really: it utilizes a bit of cleverness dubbed the "kernel trick" to keep the calculations in the native space of the data
- SVM is **really slow**: It is an order n -cubed algorithm, meaning that running SVM on a 10,000-row dataset can take up to 1000 times longer than running it on a 1000-row dataset

SVM: Example

- We'll begin by importing data on 500 stars and 500 quasars

```
##      col.ug      col.gr      col.ri      col.iz
## Min.    :-4.2274  Min.    :-2.98092  Min.    :-0.40610  Min.    :-3.69967
## 1st Qu.: 0.6613   1st Qu.: 0.09591   1st Qu.: 0.02866   1st Qu.: 0.02976
## Median : 1.1102   Median : 0.26471   Median : 0.12162   Median : 0.14411
## Mean    : 1.3196   Mean    : 0.37682   Mean    : 0.21581   Mean    : 0.18544
## 3rd Qu.: 1.7465   3rd Qu.: 0.51801   3rd Qu.: 0.25169   3rd Qu.: 0.29248
## Max.    : 6.2807   Max.    : 2.68311   Max.    : 3.39274   Max.    : 4.04392
##      mag.r      class
## Min.    :14.43   QSO :500
## 1st Qu.:17.95   STAR:500
## Median :18.75
## Mean    :18.66
## 3rd Qu.:19.47
## Max.    :24.82
```

- We will use the functions of the e1071 package below, after scaling the predictors and performing a 70-30 data split

SVM: Example With Linear Kernel

```
library(e1071)

set.seed(202) # reproducible cross-validation
tune.out <- tune(svm,class~.,data=df.train,kernel="linear",ranges=list(cost=10^seq(-2,2,by=0.2)))
cat("The estimated optimal value for C is ",as.numeric(tune.out$best.parameters),"\n")
```

```
## The estimated optimal value for C is 1.584893
```

```
resp.pred <- predict(tune.out$best.model,newdata=df.test)
mean(resp.pred!=df.test$class) ; table(resp.pred,df.test$class)
```

```
## [1] 0.1433333
```

```
##
## resp.pred QS0 STAR
##      QS0  138   29
##      STAR  14  119
```

- Compare this figure with 0.177 for logistic regression

SVM: Example With Polynomial Kernel

```
set.seed(202)
tune.out <- tune(svm,class~.,data=df.train,kernel="polynomial",
               ranges=list(cost=10^seq(2,4,by=0.5),degree=2:4))
cat("The estimated optimal values for C and degree are ",as.numeric(tune.out$best.parameters),"\n")
```

```
## The estimated optimal values for C and degree are 1000 3
```

```
resp.pred <- predict(tune.out$best.model,newdata=df.test)
(svm.poly.mcr <- mean(resp.pred!=df.test$class)) # cf. 0.177 for logistic regression
```

```
## [1] 0.1033333
```

```
table(resp.pred,df.test$class)
```

```
##
## resp.pred QSO STAR
##      QSO  130    9
##      STAR   22  139
```

- Note: if the optimal value is a boundary value (for instance, the optimal value of cost is 1 when the minimum value considered is 1), change the ranges and rerun the code
- Also, note the very sparse nature of the ranges grid: this is to ensure a relatively small computation time

SVM: Example With Radial Kernel

```
set.seed(202)
tune.out <- tune(svm,class~.,data=df.train,kernel="radial",
               ranges=list(cost=10^seq(-1,1,by=0.5),gamma=10^seq(-1,1,by=0.4)))
cat("The estimated optimal values for C and gamma are ",as.numeric(tune.out$best.parameters),"\n")
```

```
## The estimated optimal values for C and gamma are 1 0.6309573
```

```
resp.pred <- predict(tune.out$best.model,newdata=df.test)
(svm.poly.mcr <- mean(resp.pred!=df.test$class)) # cf. 0.177 for logistic regression
```

```
## [1] 0.05
```

```
table(resp.pred,df.test$class)
```

```
##
## resp.pred QSO STAR
##      QSO  142   5
##      STAR  10 143
```

- In this instance, a radial-kernel SVM model does **much** better than a logistic regression model, in terms of misclassification rate
- Here, we assume a class-separation threshold of 0.5...if you wish to create a ROC curve, you'd pass the argument `probability=TRUE` to *both* `tune()` and `predict()`, take the matrix that is output and extract the second column, and pass those numbers into `roc()`, etc.